

nmap-harvester

A supervised machine learning model designed for detecting NMAP port scanning, developed as part of a university project.

Introduction

This project aims to build a supervised machine learning model to detect real-time NMAP port scanning activities.

In many cyber-attacks, the initial step often involves port scanning using tools like NMAP. Detecting such scans can be challenging because network packets carry extensive information, and a single packet isn't enough to confirm an NMAP scan attempt.

To address this, this project proposes a machine learning-based approach for identifying TCP port scans initiated by NMAP.

Dataset Creation

Understanding TCP connections is important for building the dataset. When a TCP packet is sent over the network, it carries specific flags that facilitate the 3-Way Handshake. NMAP can manipulate these flags to evade detection while performing rapid port scans.

Key Dataset Characteristics

- **Session-Based Rows:** Instead of logging each packet individually, each row in the dataset represents a **session** (requests + responses).
- **Flag Summation:** Flags (SYN, ACK, FIN, etc.) are aggregated across the session. For example:
 - If a SYN packet is sent by NMAP and another SYN is part of the response, the SYN column will record a value of 2.

Example Dataset Row

```
start_request_time,end_request_time,start_response_time,end_response_time,src_ip,dst_ip,src_port,dst_port,duration,flag
2025-01-08 13:52:55.274814,2025-01-08 13:52:55.274814,2025-01-08 13:52:55.274874,2025-01-08 13:52:55.274874,6e-05,['172.31.0.1','172.31.0.2'],,['172.31.0.1','172.31.0.2'],,['44031','3306'],,['44031','3306'],1,1,0,1,0,0,1
```

- Sessions are grouped by `src_ip`, `dst_ip`, `src_port`, and `dst_port`. However, these grouping keys are excluded from the model's training phase.
- The `duration` feature provides valuable information for distinguishing between legitimate traffic and NMAP scans, as legitimate HTTP requests may exhibit similar flag behavior but differ in timing.
- The session window is set to **0.5 seconds** by default, as this is typically enough to capture an NMAP scan attempt.

Common NMAP Scans

```
nmap -sT 172.31.0.1 -p 0-10000 # TCP Scan
nmap -sS 172.31.0.1 -p 0-10000 # Stealth Scan
nmap -sF 172.31.0.1 -p 0-10000 # FIN Scan
nmap -sN 172.31.0.1 -p 0-10000 # NULL Scan
nmap -sX 172.31.0.1 -p 0-10000 # XMAS Scan
```

The dataset consists of: - `bad.csv`: Sessions labeled as 1 (NMAP traffic). - `good.csv`: Sessions labeled as 0 (legitimate traffic).

These datasets are combined into a single file, `merged.csv`, for training the model.

Machine Learning Model

The `XGBClassifier` was selected as the final model due to its reliable performance in key areas:

1. High MCC score (~ 0.91)
2. High accuracy score (~ 0.95)
3. Fast prediction speed ($\sim 3\text{ms}$ on average for 15,000 rows)

Model Performance Example

```
XGBClassifier (n_estimators=210): Accuracy: 0.9599, Train time: 68ms, Prediction time: 3ms, MCC: 0.919739, TP: 718, TN: 741, FN: 28, FP: 33
```

Best Classifier based on MCC

Classifier: XGBClassifier

n_estimators: 210

MCC Score: 0.919739

The speed of prediction played a significant role in choosing this model, as it allows efficient analysis of large volumes of network traffic in real time.

Training Dataset

The training dataset, `datasets/train/merged.csv`, is generated using the following approach:

1. Create an isolated Docker environment for capturing clean data:

```
docker compose up --build
```

2. Access the container:

```
docker attach traffic_generator
```

3. Run the interceptor on the host:

```
sudo python3 interceptor.py
```

- Adjust:

- `interface`: Docker network interface name
- `scanner_ip`: IP assigned to `traffic_generator`
- `output_file`: Output CSV file path
- `label`: 0 for legitimate traffic, 1 for NMAP scans

4. Run NMAP scans from the container:

```
nmap -sT 172.31.0.1 -p 0-10000
```

```
nmap -sS 172.31.0.1 -p 0-10000
```

5. Run noise traffic for legitimate requests:

```
python3 noiser.py
```

6. Merge datasets:

```
cd datasets
python3 merger.py
```

7. Train the model:

```
python3 algo_chooser.py
```

Delayed Dataset

A delayed dataset can be created by introducing delays between requests:

```
nmap -p 1-10000 --scan-delay 1s 172.31.0.1
```

You can also adjust the delay in legitimate requests by modifying `SLEEP_SECOND` in `noiser.py`.

Results:

```
Dataset loaded with 11351
records.
```

```
Dataset preprocessed
successfully.
```

	duration	SYN	ACK	FIN	RST	URG	PSH
--	----------	-----	-----	-----	-----	-----	-----

label

0	0.000060	1	1	0	1	0	0
---	----------	---	---	---	---	---	---

1

1	0.000068	1	1	0	1	0	0
---	----------	---	---	---	---	---	---

1

2	0.000062	1	1	0	1	0	0
---	----------	---	---	---	---	---	---

1

3	0.000057	1	1	0	1	0	0
---	----------	---	---	---	---	---	---

1

4	0.000074	1	1	0	1	0	0
---	----------	---	---	---	---	---	---

1

```
Dataset split into training and testing sets.
```

```
....
```

```
XGBClassifier (n_estimators=210): Accuracy: 1.0000, Train time:
28ms, Prediction time: 3ms, MCC: 1.000000, TP: 743, TN: 393, FN:
0, FP: 0
```

```
Best Classifier based on MCC
Classifier: XGBClassifier
```

n_estimators: 210
MCC Score: 1.000000

Running the Detector

To run the detector:

```
sudo python3 detector.py
```

- The detector uses `interceptor.py` to monitor session packets.
 - `injector.py` simulates normal HTTP traffic with occasional NMAP scans (~10% probability).
 - If at least **30%** of session packets are flagged as anomalies, the system will detect an ongoing NMAP attack.
-

Requirements

Install dependencies with:

```
python3 -m venv venv && source venv/bin/activate && pip install -r requirements.txt
```

Demonstration Video

<https://github.com/user-attachments/assets/f10773c6-742e-4394-913e-42beb0cc3683>

References

- [Unix Stack Exchange - Detecting NMAP Scans](#)
 - [Medium Article on NMAP Detection](#)
-

External Dependencies

- `pyshark`
- `python-nmap`